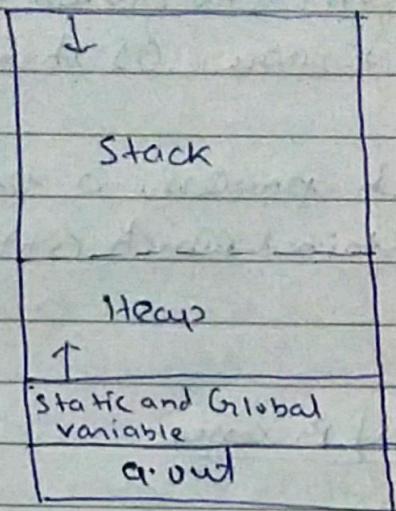


When we want to execute a program, the OS will create a process related to the program.

Stack is needed to function and recursive calls. It also contains local variables and parameters passed to the functions.
Heap is used in case of dynamic need of space.

The executable code cannot take any space beyond the process boundary.



Process id

Program counter: next instruction to be executed.

Process state

Priority: - OS process will have higher priority over user process.

General purpose registers: - When the current process is preempted during the execution then the values in the GPRS should be saved otherwise new process will replace the current GPRS values with new values. In that case, when the old process will resume its values with GPRS values will be lost. GPRS are commonly used by each process.

~~List of open devices~~

Protection :- no process should occupy other process's space. OS should ensure this.

For each process, a process control block is maintained which contains above variables.

States of Process

i) New :- The program going to be executed is called a process.

ii) Ready :- When we create a process, it is ready to run. So, it is in ready state.

When many processes are ready to run then it is called multiprogramming.

Multiprogramming is two types: with preemption (multitasking) and without preemption.

iii) Run :- On a single CPU, only one process will run at a time.

iv) Wait :- During running if the process need I/O access then the process is brought into wait state. Form

(v) termination on completion: - After finishing execution of process, the process is killed and its PCB (context) will be deleted.

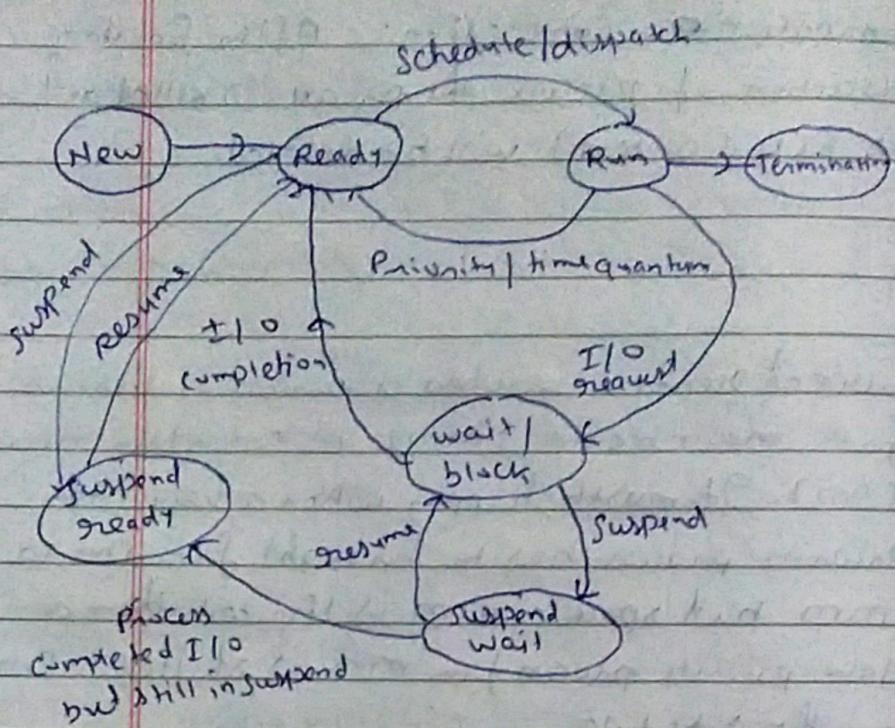
(vi)

Suspend ready: - when a process is thrown from main memory (mm) to secondary memory (sm). It mostly happens when a very high priority process has to be brought from sm to mm but space in mm is limited then a low priority process from mm is shifted to sm by the scheduler.

(vii)

Suspend wait or suspend block: - Similar to suspend ready, if in place of a process in ready state, a process in wait state is thrown to the sm. In this case, if the process finishes its wait in the sm then it can move to suspend ready so that when the process is brought back to mm, it kept in ready state and not in wait state.

Process are created, scheduled, executed & deleted



Long term Scheduler: - Gt decides how many

processes should be created or made. Gt also checks that
correct balance of CPU and I/O bound processes in the ready queue. Gt
may take time in decision but it should take efficient decisions.

Short term Scheduler: - Gt decides which

process to be shifted to run state. Gt should take the
decision quickly. e.g. if time quantum is 10ms then it has to decide within
10ms that which process should be executing next otherwise CPU will sit idle.

Middle term scheduler: - Gt decides which
process to be suspend.

Degree of multiprogramming: - max no. of
process that can be present at ready state.

Long term scheduler decides degree of multiprogramming

After choosing a process to be shifted from ready to run state, the short term scheduler ~~checks~~
calls dispatcher which whose job is to bring the chosen job from ready to run state
and shift the running process to ready in case
of Priority scheduling. This is called context
switching, i.e. changing the status of the CPU
by replacing an old process by a new process.
~~If~~ The delay in context switching should be
(a) minimum as possible.

swapping: - moving a process from mm to Sm.

LT is related to overall performance.

MT scheduler is related to swapping

STT - - - - content switching.

Consider a system with 'N' CPU & 'm' processes

Ready Ready : - min - 0

max - m

running - min - 0

max - N

block - min - 0

max - m

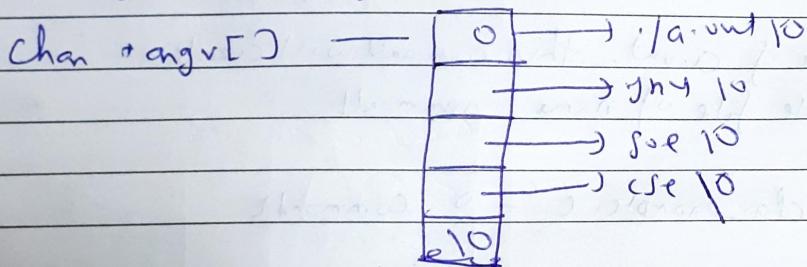
Command line arguments

M T W T F S S
Page No.:
Date: YOUVA

```
int main( int argc, char *argv[] )  
{  
    printf( " no. of arguments passed %d \n", argc );  
    printf( " parameters are \n" );  
    for( int i=0; argv[i] != NULL; i++ )  
    {  
        printf( " argv[%d] = %s, i, argv[i] );  
    }  
    return 0;  
}
```

\$./a.out Jnu sue cse

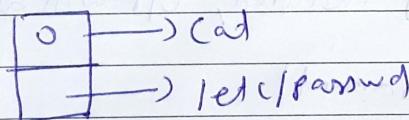
∴ argc = 6



\$ cat /etc/passwd

∴

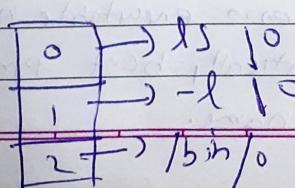
= argc = 2, argv



off

\$ ls -l bin

argc = 3, argv



AVUO

M	T	W	T	F	S	S
Page No.:						
Date:						

Here, ./a.out or cat.out are also part of argv but we can create an executable file of a program with different names e.g.

```
int main ( int argc, char *argv [] )  
{  
    if ( strcmp ( argv [0], "Command1" ) == 0 )  
        printf ("Command1 printed");  
    else  
        printf ("Command2 printed");  
}
```

O/P

```
gcc cla-example.c -o command1
```

In place of a.out, this command will create an executable file of name command1

111b,

```
gcc cla-example.c -o command2
```

Now, we have 2 executable files, command1 and command2 for the same program. But, depending upon which file is being executed, O/P will differ as per the p/m.

```
$ ./command1
```

O/P Command1 printed

```
$ ./command2
```

O/P command2 printed

so this is the reason executable name is also passed as argv[0] automatically but p/m o/p may depend on the executable name.

⇒ environment variables

\$ env

\$ set | grep -i histsize
HISTSIZE = 1000

\$ HISTSIZE = 250

\$ set | grep -i histsize
HISTSIZE = 250

This change is temporary change. To change it permanently, change in .bashrc file for a particular user and /etc/profile for all users of the system.

\$ man environ

~~extern char **environ;~~

int main()

{

 printf("env variables are: ">,

 for (int i = 0; environ[i] != NULL; i++)

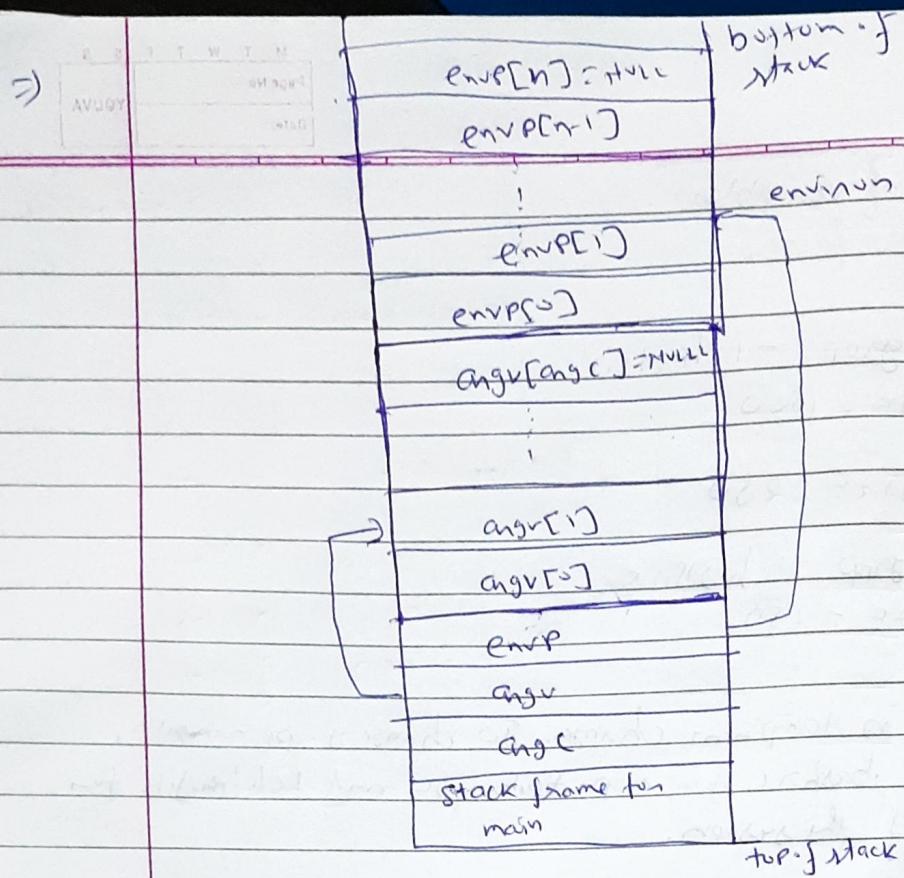
 printf("%s\n", environ[i]);

}

* It will print all environment variables. All ~~/* env comment~~ printed.

putenv & setenv functions can be used to change the value of env. variable

getenv can be used to access a particular env. variable.



Layout of process stack

⇒ Function Stack frame
of a function containing

- (a) local variables of the function
- (b) argument passed to the function
- (c) return address
- (d) base pointer

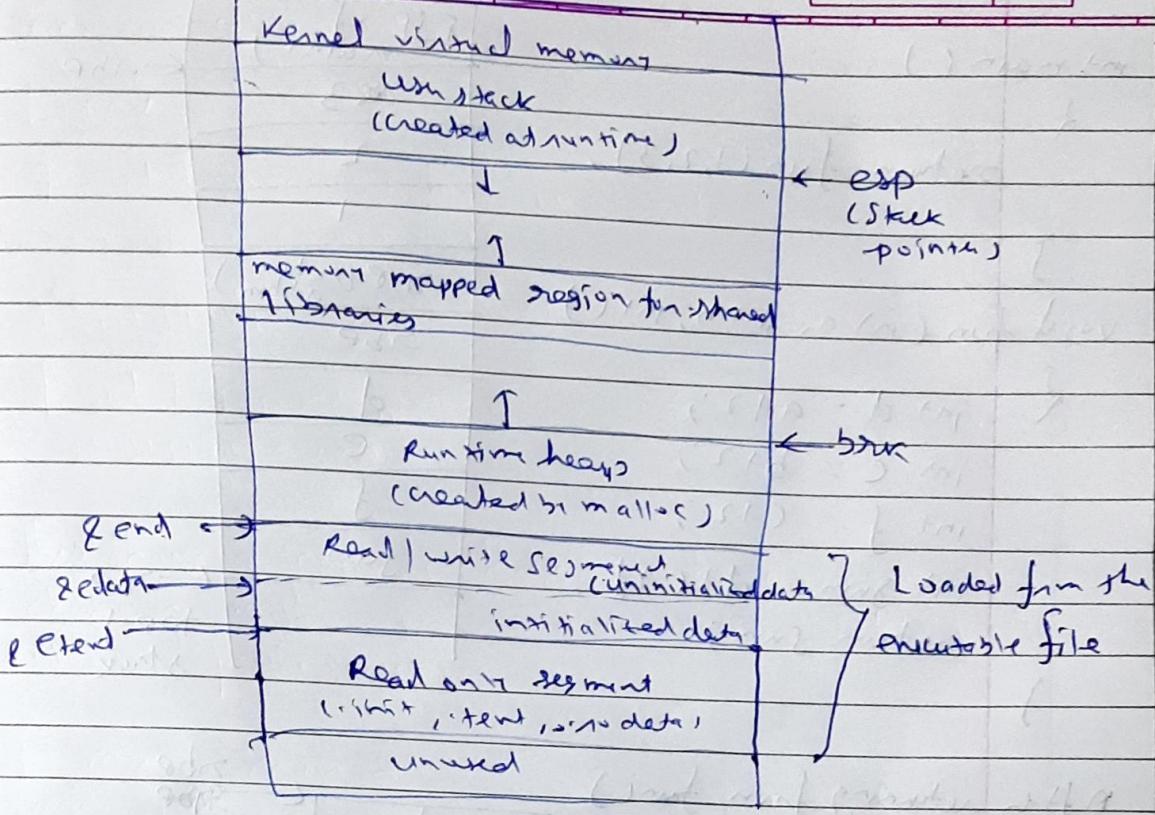
argc, argv, envp
FSF of main()

← sfp
← top of stack

NSP & DOP are base pointers.

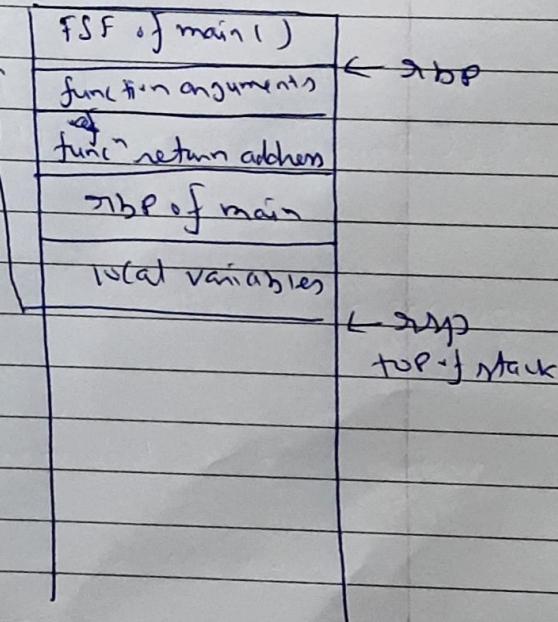
sfp points to beginning of
FSF of the function.

NSP points to top of the stack.



In main() function is calling fun() function

- Arguments are pushed in stack in reverse order
- The contents of rbp containing starting address of main stack frame is saved on stack for late use and rbp is moved to where rbp is pointing to create new stack frame pointer of function fun()



```

int main()
{
    return fun(1, 2, 3);
}
    
```

FSF of
main()

*3

2

1

visit fun(int a, int b, int c)

top of stack (return address)

{ int d = a + 2;

ebp

int e = b + 2;

d

int f = c + 2;

e

int sum = d + e + f;

f

return sum;

sum

}

top of stack

After returning from fun()

FSF of
main()

ebp

top of stack

Buffer overflow - It is a bug in a program which occurs when more data is written to a block of memory than it can handle.

```
int main (int argc, char *argv[])
{
    display (argv[1]);
}
```

```
void display (char *str)
```

```
{
    char buff[10];
    strcpy (buff, str);
    printf ("data is %s\n", buff);
}
```

FSF } main()

*Str

rip

rbp

buff

Process

\$ ps -ef | grep -i kernel

Kernel keeps track of ^{mgt} pages occupied by different processes.

In the top of process stack, ^(containing) Kernel Virtual memory address space is there.

The kernel address space is divided into 2 parts, part 1 contains kernel code & its data structures which are same for all processes whereas part 2 contains process specific data structures which are different for every process.

e.g. state of a process, resources occupied by the process, etc.

\$ uname -a

\$ cd /usr/src/linux-headers-5.4.0-90-generic/include/linux

\$ less fs.h

:/struct file {

to check the structure of a file.

\$ less sched.h

(to see structure related to process management, which is PCB)

:/struct task_struct {

explore the different variables declared in this task structure.

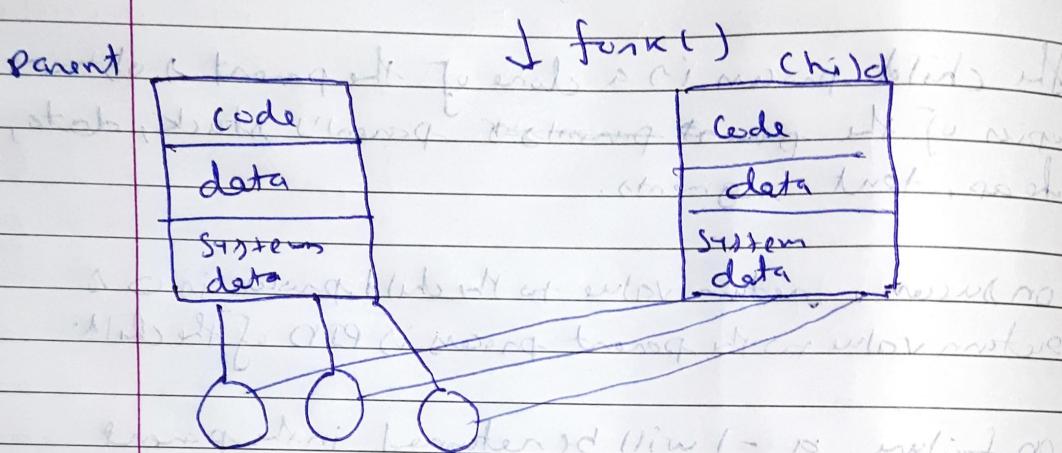
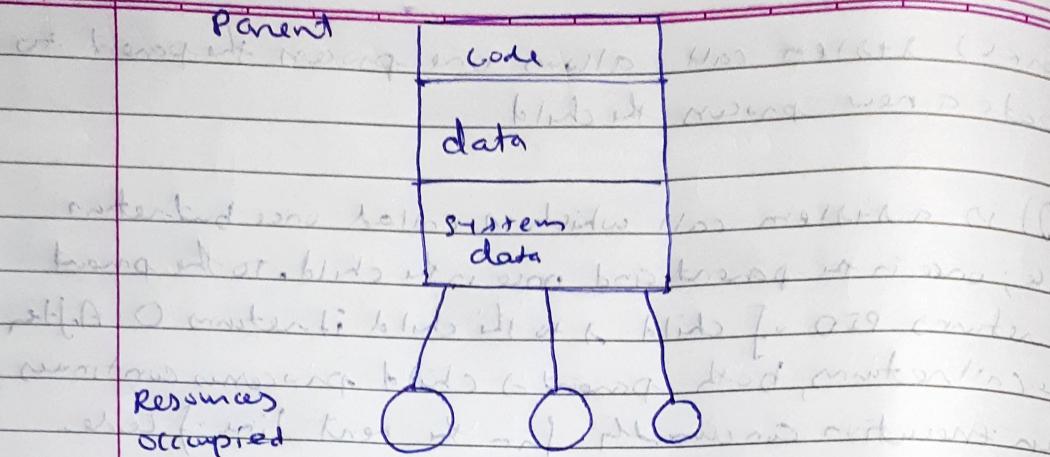
④ \$ cd /proc check the files related to a process in /proc directory.

S	S	T	W	T	F	S
A	V	C				
U	N	E	R	P	E	S

M T W T F
Page No.:
Date:

- ⇒ getpid() and getppid() can be used to get the pid of a process and its parent process.
- ⇒ In a shell, the PID of the shell can be get through the environment variable \$~~\$~~ \$\$ and its parent by &PPID.
- ⇒ Swapper or Scheduler processes (schedules in a system) have PID as 0. The program file for swapper in /proc directory.
- ⇒ Init on Systemd has a PID 1.
- ⇒ page daemon on Kthreadd has PID 2.
- ⇒ It supports the paging of virtual memory system.

- fork() system call allows one process the parent to create a new process, the child.
- It is a system call which is called once but returns twice; once in the parent and once in the child. To the parent it returns PID of child & to the child it returns 0. After the call returns, both parent & child processes continues their execution concurrently from the next line of code.
- The child process is a clone of the parent & obtains copies of the parent's stack, data, heap, text segments.
- On success, return value to the child process is 0 & return value to the parent process is PID of the child.
- On failure, a -1 will be returned in the parent content, no child process created.
- After a fork(), it is indeterminate which process, parent or child will execute. → In most of the cases, parent will execute but parent's state is already active in the CPU → its memory mgmt. info is already cached.
- fork1.c - simple fork program
- fork2.c - two child & parent runs concurrently
two a.out can be seen by PS - a command.



Accessing parent algorithm - > parent

Accessing own knows about its own - > parent

Accesses & local access and local environment

2

Orphan Process

→ If a parent has terminated before reaping its children and the child process is still running, then these children are called orphans.

→ They are adopted by init or systemd process to see whether the process has adopted by init or not
\$ ps -l

How to check the PPID of a pid
now, ~~now~~ ~~for~~

\$ ls -l /proc/~~pid~~/one

It will give the name of the process having ~~0~~ process id as ppid.

Zombie Process

→ processes which have terminated but their parents have not collected their exit status. It has not reaped them and called zombies or defunct. So, a parent must reap its child.

→ When a process terminates but still is holding system resources like PCB & various tables maintained by OS. It is half-alive & half dead but it is holding resources like memory but it is never scheduled on the CPU.

→ Zombies can't kill by a signal, not even with the SIGKILL. The only way to remove them from the system is to kill their parent, at which time they become orphan or adopted by init or systemd.

Let's do zombie.c, child will become zombie. To check, go to another terminal & use command PS -a. It will show 'defunct' process.

\$ kill %1 of zombie (not killed)

\$ kill -9 %1 of zombie (not killed)

\$ K11 - 9 pid = 5 parent of zombie

replaced after capture and became infected long after
 now zombie presumed parent will be init.
 Red collar, camera will be set up to help
 and give better surveillance

During last time I'd brought an adult
 female and male and among all others were
 (1) radio

turns to 0799 and determine

adult male and female and add more to
 existing group 1B - 1B

which was saved among all from previous notes

radio and male and female and female among
 no male because ten and a female has radio. Between them
 did they have enough time to get treated or evidence being
 described in visited station because a radio

was 1.0 and station silent during a 028 and was
 never established and had fresh & swollen
 08 above both sides were affected areas 201

with new skin was too large and (no time) evidence
 that it was not just one of many other things
 though it was not good for the body to have a
 bacteria on their

and not good for health so sides open
 (0400 am 19) having new to ferment others of
 many bacteria and

(0411 am 19) signs of 0469 the
 character was stable R-11

Wait to system call

Parent process can monitor the execution or termination of its child process through wait() system call.

Pid = wait (int *status)

- > the process that calls the wait() system call gets blocked till any one of its child terminates
- > the child process returns its terminal status using the exit() system call & that integer value is received by the Parent inside the status argument
- > On success, the wait() system call returns PID of the terminated child & in case of error, returning -1.
- > If a process wants to for terminate of all its children, then while(wait(null) > 0);
- > A process can end in four ways
 - (a) Success / failure
 - (b) Killed by a signal
 - (c) Stopped by a signal
 - (d) continued by a signal

FNC functions

- A process may overwrite itself with another executable image. When a process calls one of the six `exec()` func's, it is completely replaced by the new program → the new program starts executing its main function.
- There are five library func's of exec family → all are layered on top of `execve()` system call. Each of these functions provides a different interface to the same functionality.
- There is no return after a successful exec call. The `exec()` functions return only if an error has occurred. The return value is -1 & ~~errno~~ `errno` is set to indicate the error.

exec functions

→ A process may overwrite itself with another executable image. When a process calls one of the six exec() funcs, it is completely replaced by the new program → the new program starts executing its main function.

→ There are five library funcs of exec family → all are layered on top of execve() system call. Each of these functions provides a different interface to the same functionality.

→ There is no return after a successful exec call. The exec() functions return only if an error has occurred. The return value is -1 → errno errno is set to indicate the error.

```
int exec (const char *pathname, const char *args, ..., (char *) 0);
int execvp (const char *filename, const char *args, ..., (char *) 0);
int execle (const char *pathname, const char *args, ..., (char *) 0,
            char *const envp[]);
```

→ The 1st argument to this family of exec() calls, is the name of the executable, which on success will overwrite the address space of the calling process with a new p/m from the 2nd storage.

→ The 'l' after the exec means that command line arguments to the new p/m will be passed as a comma separated list of strings with a '\0' character at the end.

→ The 'p' stands for path. It means that the p/m specified as the 1st argument should be searched in all directories listed in the PATH variable. However, using absolute path

to pjm is more secure than setting on PATH variable, which can be more easily altered by malicious user.

→ The 'e' stands for environment. It means that after the command line arguments, the pjm should pass an array of pointers to null terminated strings, specifying the new environment of the pjm to be executed. Otherwise the caller environment will be used.

execv family

```

int execv (const char * pathname, char * const argv[]);  

int execvp (const char * filename, char * const argv[]);  

int execvp  

int execvr (const char *file + pathname, char * const argv[],  

            char * const envp[]);
    
```

here '*' means the list of argument will not be passed as comma separated list but as an array of pointers.