

message queues

→ mqs can be used to pass messages b/w related or unrelated processes executing on same machine. message queues are somewhat like pipes & fifos but differ in some important aspects.

- (1) message boundaries are preserved so that reader and writer communicate in units of message.
- (2) ~~message queues are kernel persistent~~.

→ as mq can be thought of as a linked list of messages in the kernel space. processes with adequate permission can put messages onto the queue & processes with adequate permission can remove messages from the queue.

→ In pipes, a process read or write streams of bytes, while in mq a process read or write a complete delimited message. It is not possible to read a partial message leaving rest behind in IPC object.

→ In pipe, a write makes no sense unless a reader also exists. In mq, there is no requirement that some reader must be waiting before a process writes a message to the queue.

→ mqs are priority driven. Queue always remains sorted with the ~~oldest~~ oldest message of the highest priority at front.

→ A process

Creating or opening a message queue
 int msgget (key_t key, int msgflag)

- To create a brand new MQ, or to get the identifier of an existing queue, we use the msgget() system call which on success returns a unique MQ identifier.
- If a MQ associated with the 1st argument, key already exists, the call returns the identifier of the existing MQ otherwise it creates a new one.
- We can use IPC_PRIVATE as a 1st argument: A parent process creates MQ prior to performing fork(); & the child inherits the returned MQ identifiers so that parent's child can communicate using this MQ. For unrelated process, we can use IPC_PRIVATE but in that case the newer process has to write the returned MQ identifier in a file that can be read by the other process.
- The 2nd argument msgflag is normally IPC_CREAT | O_CREAT
- Instead of IPC_PRIVATE constant as 1st argument, processes can use the ftsk() library call with same arguments to generate a unique key which can be used by msgget() function.

```
int msgsnd (int msgid, const void * msgp, size_t msgsz,
           int msgflg);
```

→ 1st argument is used to send a message, which is the to the msg identifier 2nd argument, which is the msg identification.

→ 2nd argument is a pointer to a structure of type msgbuf having following 2 fields:

struct msgbuf

```
{ long mtype; // used to retrieve a message by type
  char mtext[512];
};
```

→ 3rd argument msgsz is the size of mtext field in the structure msgbuf.

→ 4th argument msgflg can be 0 or IPC_NOWAIT . i.e. if msgid has reached its max. no. of messages, in that case, the sender will wait until the msg queue does not have enough space for new msg. If it is 1 means the sender process will immediately return and will not add the msg in msg queue.

```
int msgrcv (int msgid, void * msgp, size_t maxmsgsz,
            long msgtype, int msgflag)
```

→ The msgrcv() system call reads & removes a msg from msg queue identified by its 1st argument msgid → copies its contents into the buffer provided by its 2nd argument msgp

→ 3rd argument maxmsgsz, specifies the maximum space available in the mtext field.

→ The 4th argument msgtype is used to specify as to which msg is to be removed & returned to the called caller based.

This is achieved by specifying msgtype field of the struct msgbuf.

msgtype == 0

1st msg from msg queue removed & returned

msgtype > 0

1st msg from queue whom msgtype field equals to msgtype is removed and returned to calling process

msgtype < 0

1st msg w/ the lowest msgtype field less than or equal to absolute value of msgtype is removed & returned.

msg position in queue

msgtype

msg body

1

300

-

2

100

-

3

200

-

4

400

-

5

100

-

If msgtype = 0, i.e. msgrcv(id, &msg, maxmsg, 0, 0)

then receiver will receive msg in following order:

1 (msgtype = 300)

2 (msgtype = 100)

3 (msgtype = 200)

4 (msgtype = 400)

5 (msgtype = 100)

If msgrcv(id, &msg, maxmsg, 100, 0)

the receiver will receive msg in following order

2 (msgtype = 100)

5 (msgtype = 100)

if `msgrecv(id, &msg, msgsize, -3w, 0)`
 the receiver will receive msg on following
 order:

2 (`mtypes = 1w`)

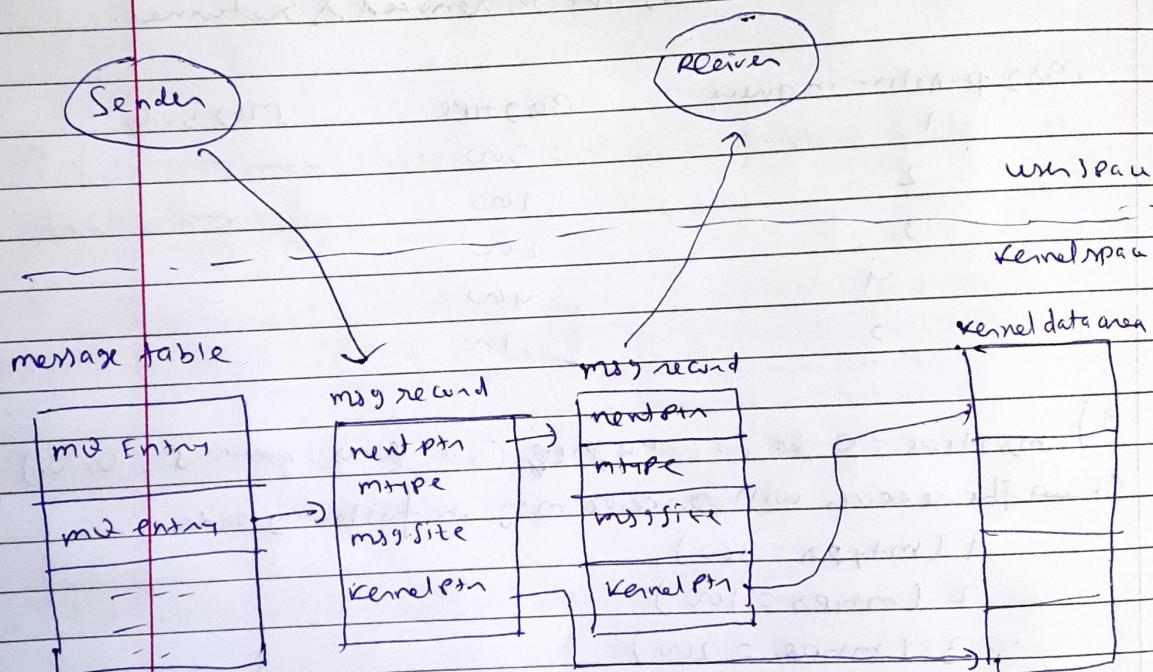
5 (`mtypes = 1w`)

3 (`mtypes = 2w`)

1 (`mtypes = 3w`)

next call of `msgrecv` will block the process.

Here, notice that one `msgrecv` call will retrieve only one msg. To retrieve multiple messages like above, multiple `msgrecv` calls are needed.



M	T	W	T	F	S	S
Page No.:		Date:	YOUVA			

→ to see the messages of messages M & C can be created by the kernel
\$ cat /proc/sys/kernel/msgmni
32000

→ to see messages that can be written in a msgq { amq.
\$ cat /proc/sys/kernel/msgmax
8192

\$ ipcs -q

\$ gcc sender.c -o sender

\$./sender

\$ ipcs -q